

# OutOfMemoryError: What is the cost of Java objects

Jean-Philippe Bempel  
Senior Software Engineer  
 @jpbempel



# **java.lang.Object**



# Anatomy of java.lang.Object header

Fields	32 bits	64 bits	64 bits CompressedOops
mark word	4 bytes	8 bytes	8 bytes
klass pointer	4 bytes	8 bytes	4 bytes
Total	8 bytes	16 bytes	12 bytes (but 16 with padding/alignment)



## Mark Word

Bitfields			Tag	State
HashCode	Age	0	01	Unlocked
Lock record address			00	Light-weight locked
Monitor address			10	Heavy-weight locked
Forwarding address, etc.			11	Marked for GC
Thread ID	Age	1	01	Biased / biasable




# Compressed Oops



# Compressed Oops

- 32 bits pointer can address 4GB
- Memory addresses is aligned (4 or 8)
- Objects reside only on address 8 multiple
- Last 3 bits are not effectively used
- Use this 3 bits to increase addressable space





## Example:

Address 88 in binary:

101 1000

will be encoded as (3 bits shift)

1011

Allows to encode 8 times more Object addresses than with raw 32 bits addressing.  
=> Maximum addressable is then 32GB.

CompressedOops saves ~20-30% memory

Activated by default on 64bits JVM

JVM Options: -XX:+UseCompressedOops



# Padding





# Padding

Heap Dump analysis with profilers can give you a good estimation

Real size depends on 32bits/64bits/CompressedOops

A more precise way: Java Object Layout (JOL)

```
java -XX:+UseCompressedOops -jar jol-internals.jar java.lang.Object
```

```
Running 64-bit HotSpot VM.
```

```
Using compressed references with 3-bit shift.
```

```
Objects are 8 bytes aligned.
```

```
Field sizes by type: 4, 1, 1, 2, 2, 4, 4, 8, 8 [bytes]
```

```
Array element sizes: 4, 1, 1, 2, 2, 4, 4, 8, 8 [bytes]
```

```
java.lang.Object object internals:
```

OFFSET	SIZE	TYPE	DESCRIPTION	VALUE
0	4		(object header)	01 00 00 00 (0000 0001 0000 0000 0000 0000 0000 0000)
0000)				
4	4		(object header)	00 00 00 00 (0000 0000 0000 0000 0000 0000 0000 0000)
0000)				
8	4		(object header)	6d 05 88 df (0110 1101 0000 0101 1000 1000 1101 1111)
1111)				
12	4		(loss due to the next object alignment)	

```
Instance size: 16 bytes (estimated, add this JAR via -javaagent: to get accurate result)
```

```
Space losses: 0 bytes internal + 4 bytes external = 4 bytes total
```



# Padding

With 8 bytes field alignment padding occurs

What is the real size of this class:

```
class Data
{
    long l;
    boolean b;
    int i;
    char c;
    String str;
}
```



# Padding

## Class Data:

```
java -XX:+UseCompressedOops -classpath .;jol-internals.jar org.openjdk.jol.MainObjectInternals Data
Running 64-bit HotSpot VM.
Using compressed references with 3-bit shift.
Objects are 8 bytes aligned.
Field sizes by type: 4, 1, 1, 2, 2, 4, 4, 8, 8 [bytes]
Array element sizes: 4, 1, 1, 2, 2, 4, 4, 8, 8 [bytes]
```

Data object internals:

OFFSET	SIZE	TYPE	DESCRIPTION	VALUE
0	4		(object header)	01 00 00 00 (0000 0001 0000 0000 0000 0000
0000 0000)				
4	4		(object header)	00 00 00 00 (0000 0000 0000 0000 0000 0000
0000 0000)				
8	4		(object header)	b6 8c 91 df (1011 0110 1000 1100 1001 0001
1101 1111)				
12	4	int	Data.i	0
16	8	long	Data.l	0
24	2	char	Data.c	
26	1	boolean	Data.b	false
27	1		(alignment/padding gap)	N/A
28	4	String	Data.str	null

```
Instance size: 32 bytes (estimated, add this JAR via -javaagent: to get accurate result)
Space losses: 1 bytes internal + 0 bytes external = 1 bytes total
```



# Structure sizes



# Arrays

Arrays have additional int for size

header for 64bits + CompressedOops: 16 bytes

byte[32]	=>	header	+	1*32	bytes	=	48	bytes
short[32]	=>	header	+	2*32	bytes	=	80	bytes
char[32]	=>	header	+	2*32	bytes	=	80	bytes
int[32]	=>	header	+	4*32	bytes	=	144	bytes
long[32]	=>	header	+	8*32	bytes	=	272	bytes
double[32]	=>	header	+	8*32	bytes	=	272	bytes
Object[32]	=>	header	+	RefSize*32	bytes	=	144	bytes



# String class fields

1.6.0\_45 (32 bytes)

- char[] value
- int hash
- int count
- int offset

1.7.0\_55 (24 bytes)

- char[] value
- int hash
- int hash32

1.8.0\_25 (24 bytes)

- char[] value
- int hash


11 (24 bytes)

- byte[] value
- int hash
- byte coder



# java.lang.String

Class String + char[]



example for string foo

Class String Size + char array header + 2\*3 bytes = 24 + 16 + 6 = 46 bytes

overhead = 1433%

for string of 100 characters:

Class String Size + char array header + 2\*100 bytes = 24 + 16 + 200 = 240 bytes

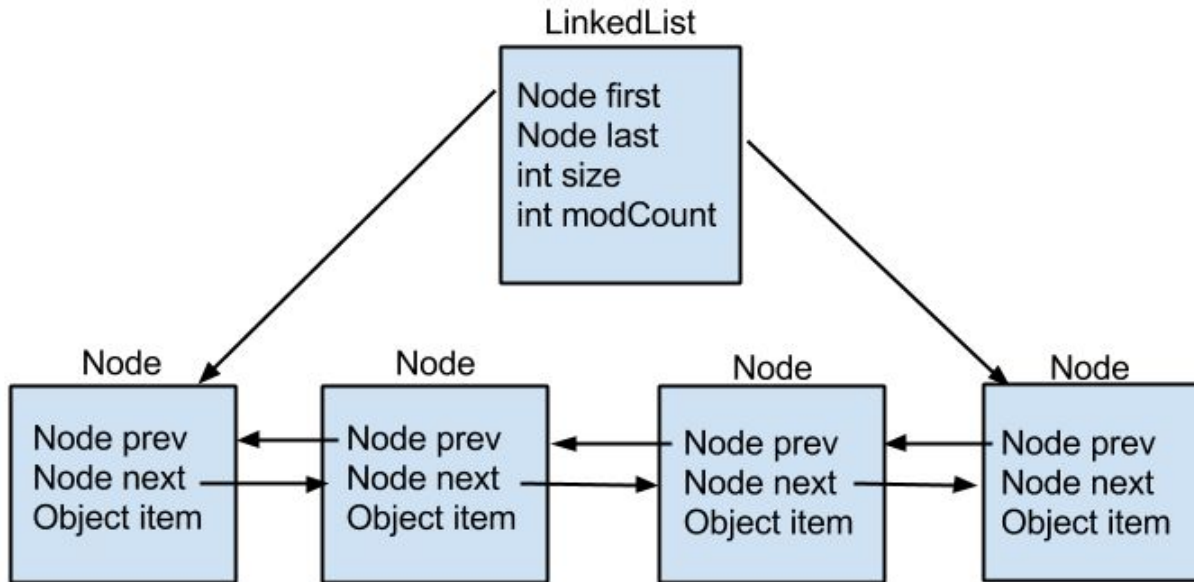
overhead = 140%

String.substring()

- <1.7.0\_06 => shared char[]
- >= 1.7.0\_06 => make copy of char[]



# java.util.LinkedList



Class `LinkedList`: 32 bytes


Class `Node`: 24 bytes

Example for 100 elements:  $32 + 100 \times 24 = 2432$  bytes





# java.util.ArrayList



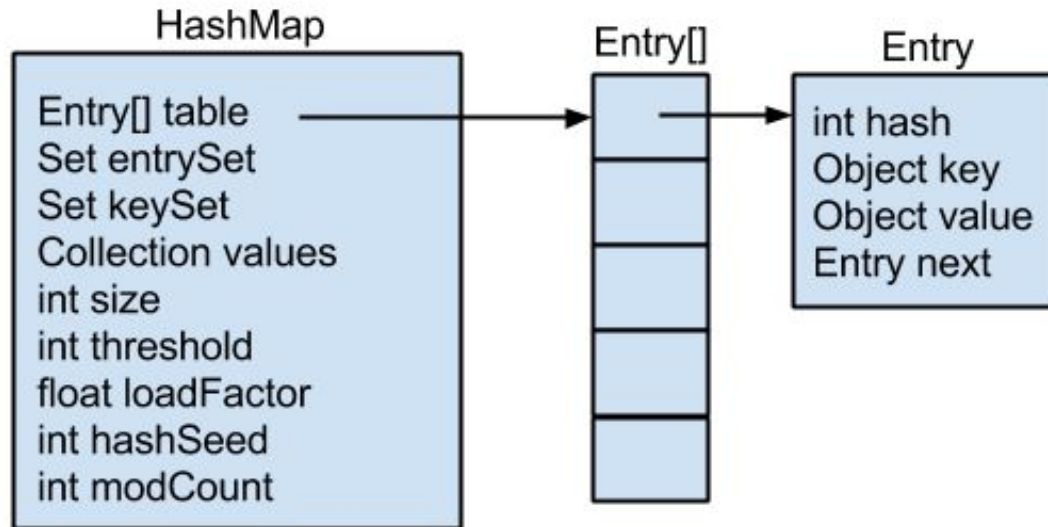
Class `ArrayList`: 24 bytes

`Object[]` :  $16 + n \cdot 4$  bytes

Example for 100 elements:  $24 + 16 + 100 \cdot 4 = 440$  bytes



# java.util.HashMap



class `HashMap`: 48 bytes

`Entry[]`:  $16 + n \cdot 4$  bytes

class `Entry`: 32 bytes

Example for 100 key/value pairs:

$$48 + 16 + 256 \cdot 4 + 100 \cdot 32 = 4288 \text{ bytes}$$



# java.util.HashMap

- `entrySet()` called => add `EntrySet` instance (16 bytes)
- `keySet` called => add `KeySet` instance (16 bytes)
- `values()` called => add `Values` instance (16 bytes)

For those inner classes this is object header + outer ref

`java.util.HashMap.Values` object internals:

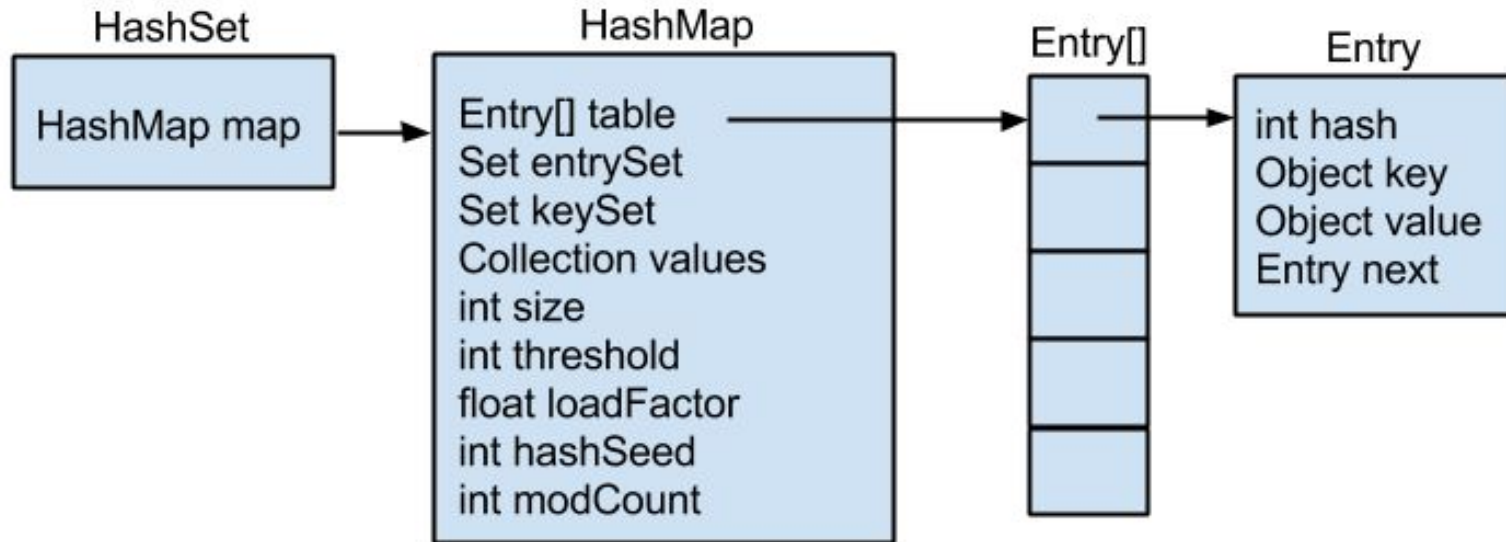
OFFSET	SIZE	TYPE	DESCRIPTION	VALUE
0	12		(object header)	N/A
12	4	<code>HashMap</code>	<code>Values.this\$0</code>	N/A

Instance size: 16 bytes (estimated, the sample instance is not available)

Space losses: 0 bytes internal + 0 bytes external = 0 bytes total



# java.util.HashSet

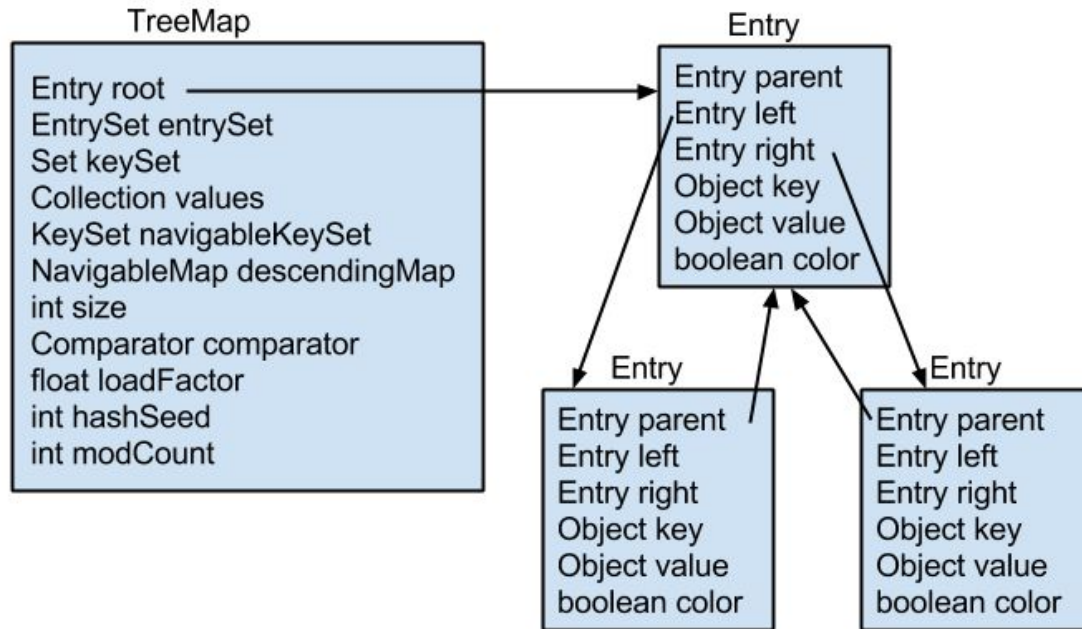


class HashSet: 16 bytes

example for 100 elements: 16 + HashMap cost = 4304 bytes



# java.util.TreeMap



class TreeMap: 48 bytes

class Entry: 40 bytes (double penalty)

example for 100 elements:

$48 + 100 \times 48 = 4848$  bytes

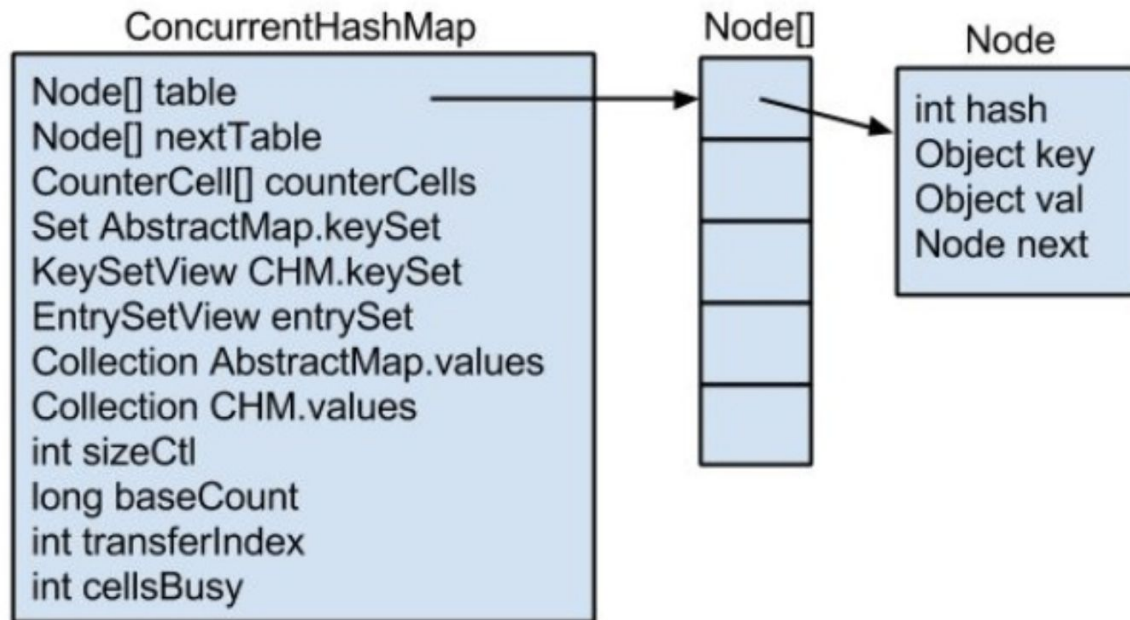


# java.util.TreeMap

- `entrySet()` called => add `EntrySet` instance (16 bytes)
- `keySet` called => add `KeySet` instance (16 bytes)
- `values()` called => add `Values` instance (16 bytes)
- `navigableKeySet` called => add `KeySet` instance (16 bytes)
- `descendingMap` called => add `DescendingSubMap` instance (56 bytes)



# java.util.concurrent.ConcurrentHashMap



class `ConcurrentHashMap`: 64 bytes

class `Node`: 32 bytes

example for 100 elements:

$64 + 16 + 256 \cdot 4 + 100 \cdot 32 = 4304$  bytes



# Summary

Collections	Overhead for 100
ArrayList	440
LinkedList	2432
TreeMap	4048
HashMap	4288
HashSet	4304
ConcurrentHashMap	4304





# Troubleshooting



# Class Histogram

jmap -histo <pid>

num	#instances	#bytes	class name
1:	73903	5978960	[C
2:	4309	4085624	[I
3:	2790	1730968	[B
4:	49641	1191384	java.lang.String
5:	22080	706560	java.util.HashMap\$Node
6:	5420	571328	java.lang.Class
7:	7408	431080	[Ljava.lang.Object;
8:	2908	331584	[Ljava.util.HashMap\$Node;
9:	1339	289224	sun.java2d.SunGraphics2D
10:	2882	253616	java.lang.reflect.Method
11:	6586	227248	[Ljava.lang.Class;
12:	1632	223768	[Ljava.lang.String;
13:	4973	198920	java.security.AccessControlContext
14:	3880	186240	java.util.HashMap
15:	4890	156480	java.util.Hashtable\$Entry

JVM Option: -XX:+PrintClassHistogramBeforeFullGC  
Includes class histogram in GC logs



# Heap Dump

To generate it:

- `jmap -dump:live,format=b,file=heap.hprof <pid>`
- JVM Options:
  - `-XX:+HeapDumpOnOutOfMemoryError`
  - `-XX:HeapDumpPath=../logs`
- JMX:  
`com.sun.management:type=HotSpotDiagnostic.dumpHeap(filename, liveonly)`

To exploit heap dump:

- visualVM
- JMC + JOverflow
- YourKit
- Eclipse MAT



VisualVM 1.3.6

FileApplicationsViewToolsWindowHelp

Applications X

Local

- VisualVM
- IntelliJ IDEA (pid 4936)
- com.ullink.testtools.fixsender.FixSenr
- com.ullink.testtools.fixsender.FixSenr
- Eclipse (pid 14704)
- com.ullink.ulbridge.Main (pid 14588)
- [heapdump] 15:54:57
- ul-launcher-0.1.jar (pid 9052)

Remote

- uranus-2
- archi-srv
  - archi-srv:18001 (pid 14851)
  - archi-srv:9091 (pid 30350)

Logfiles

Snapshots

Start Page xcom.ullink.ulbridge.Main (pid 14588) x

Overview

Monitor

Threads

Sampler

Profiler

MBeans

Buffer Pools

SA Plugin

Memory Pools

Visual GC

Tracer

[heapdump] 15:54:57 x

com.ullink.ulbridge.Main (pid 14588)

Heap Dump

Summary

Classes

Instances

OQL Console

java.util.HashMap\$EntryInstances: 104 053 | Instance size: 44 | Total size: 4 578 332 | Compute Retained Sizes

Instances

Fields

References

Instance
<500 instances>
<500 instances>
<500 instances>
<500 instances>
<500 instances>
<500 instances>
<500 instances>
<500 instances>
<500 instances>
#41501
#41502
#41503
#41504
#41505
#41506
#41507
#41508
#41509

Field	Type	Value
this	HashMap\$Entry	#41501
hash	int	2035558222
next	<object>	null
value	String	#66464
key	String	#65492
<classLoader>	<object>	null

Field	Type	Value
this	HashMap\$Entry	#41501
next	HashMap\$Entry	#41500
[14]	HashMap\$Entry[]	#7560 (16 items)
table	HashMap	#5586
h21	BiDirectionMap	#1816
mPosAmtType	CMSNormalizer	#3

Array type | Object type | Primitive type | Static field | GC Root | Loop



**Main-2015-01-13.snapshot [C:\Users\jpbempel\Snapshots] - YourKit Java Profiler 2013 build 13068 - 64-bit**

File View Memory CPU Settings Tools Help

Support period expired on November 20, 2014. To renew, [contact YourKit](#)

Exceptions \* Performance Charts \* Probes \* Inspections \* Summary \*

CPU \* Threads \* Memory \* GC Roots -> Instances of class 'java.util.HashMap\$Entry' \* Garbage Collection \*

**Paths**

Show 500 shortest paths ☒ Ignore soft/weak/finalizer references Ignore Selected Reference Undo Last Ignored Ref Undo All Ignored Refs (0)

Name	Retained Size	Shallow Size
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
[0] of java.util.HashMap\$Entry[32]	880	144
table of java.util.HashMap size = 23	928	48
h12 of com.ullink.ulbridge2.ulcms.runtime.BiDirectionMap	2 112	32
mINSTATTRIBTYPE of com.ullink.ulbridge2.ulcms.runtime.CMSNormalizer	832 240	21 152
normalizerInst of com.ullink.ulbridge2.ulcms.runtime.misc.MTNormalizer	24	24
normalizer of com.ullink.ulbridge2.plugins.FIXCMSPlugin [Stack Local]	75 432	896
<local variable: this> of java.lang.Thread [JNI Global, Stack Local, Thread] "CLICMS1"	2 624	104
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32
java.util.HashMap\$Entry	32	32

**Legend**

- Regular object
- Array
- Class
- Class loader

Working with Paths



# War stories



# Overhead structures example

```
251614.001: [GC 251614.001: [ParNew (promotion failed): 943744K->940573K(943744K), 0.9248570
secs]251614.926: [Class Histogram:
```

num	#instances	#bytes	class name
1:	421751083	20244051984	java.util.concurrent.ConcurrentHashMap\$HashEntry
2:	23327688	12130397760	com.ullink.ulmonitoring.api.model.SimpleBMOOrder
3:	23428922	6694953456	[Ljava.util.concurrent.ConcurrentHashMap\$HashEntry;
4:	11574497	5007170576	[C
5:	178973	2246627208	[B
6:	23357783	1681760376	java.util.concurrent.ConcurrentHashMap
7:	26115162	1253527776	java.util.HashMap\$Entry
8:	23386843	1122568464	java.util.concurrent.locks.ReentrantLock\$NonfairSync
9:	23381591	1122316368	java.util.concurrent.ConcurrentHashMap\$Segment
10:	27363594	1094543760	java.lang.String
11:	22350856	894034240	com.ullink.ulmonitoring.api.model.IndentedHelper
12:	22350849	894033960	com.ullink.ulmonitoring.api.model.OrderLine
13:	23357784	747639528	[Ljava.util.concurrent.ConcurrentHashMap\$Segment;
14:	22350792	715225344	com.ullink.ulmonitoring.extension.view.HierarchicalOrderLine
15:	23356168	560548032	com.ullink.ulmonitoring.api.model.property.PropertyHolderImpl
...			
817:	40	4480	
			com.ullink.ulmonitoring.extension.view.HierarchicalViewDataContainer



# Client class (104 bytes)

```
com.ullink.oms.model.Client object internals:
  OFFSET  SIZE          TYPE DESCRIPTION             VALUE
    0     4               (object header)      01 00 00 00 (0000 0001 0000 0000 0000 0000
0000 0000)
    4     4               (object header)      00 00 00 00 (0000 0000 0000 0000 0000 0000
0000 0000)
    8     4               (object header)      3b de 91 df (0011 1011 1101 1110 1001 0001
1101 1111)
   12     4             String OdisysObject.id         null
   16     4             String Client.type             null
   20     4           Capacity Client.capacity          null
   24     4             String Client.description       null
   28     4             String Client.currency          null
   32     4             String Client.parentClientId    null
   36     4           Collection Client.contact          null
   40     4             String Client.data              null
   44     4           Contact Client.office              null
   48     4             String Client.originCountry     null
   52     4           Boolean Client.allowUnknownAccount null
   56     4             String Client.connectionDetails  null
   60     4             String Client.backoffice         null
   64     4             Policy Client.policy             null
   68     4 ConfirmationSchedule Client.confirmationSchedule null
   72     4           Collection Client.stampExemptCountries null
   76     4             Map Client.customProperties       null
   80     4           Delivery Client.delivery           null
   84     4           Warehousing Client.warehousing      null
   88     4             Map Client.externalIds            null
   92     4             String Client.feeId               null
   96     4           Boolean Client.active              null
  100     4               (loss due to the next object alignment)
Instance size: 104 bytes (estimated, add this JAR via -javaagent: to get accurate result)
Space losses: 0 bytes internal + 4 bytes external = 4 bytes total
```





# Event user.update

num	#instances	#bytes	class name
-----			
1:	116925840	18595529736	[C
2:	100508918	10452927472	com.ullink.oms.model.Client
3:	116908590	2805806160	java.lang.String
4:	39690	773638968	[B
5:	77391	596989120	[Ljava.lang.Object;
6:	6116590	195730880	java.util.HashMap\$Entry
7:	2370792	149697448	[Ljava.util.HashMap\$Entry;
8:	2406042	115490016	java.util.HashMap
9:	149661	20530168	<constMethodKlass>
10:	149661	19166816	<methodKlass>

## Event user.update

1 user with 32 560 Clients

action with User with 32 560 Clients

3796 events = 245M Client instances => 25GB + Strings



# Instrument class

- 62 ref fields
- String, BigDecimal, Boolean, Double, Collection, Custom Classes (Alternateld (56B), Time (80B), Enums (24B)...)

Size: 264 bytes



# Instrument Referential: 3M

Only Instrument instances: 792MB

Add 8 chars for instrumentId: +168MB ( $56 \times 3M$ )

Add 4 chars for symbol: + 312MB ( $48 \times 3M + 56 \times 3M$ )

Add settlement date: + 408MB ( $56 \times 3M + 80 \times 3M$ )

Add 3 chars for currency: + 138MB ( $46 \times 3M$ )

Total 1818MB

Market Data Representation:

for each instrument : `Map<String, MDRecord>`

MDStringRecord: str + converted String (24B)



# Solutions



# Solutions

- Flyweight/internalizers
- ArrayList vs LinkedList
- HashMap => OpenAddressingHashMap



# Measure, don't guess!

Kirk Pepperdine & Jack Shirazi



**Measure, don't premature!**



# References

[Java Object Layout](#)

[What Heap Dumps Are Lying To You About](#)

[From Java code to Java heap](#)

[Building Memory-efficient Java Applications: Practices and Challenges](#)





# Thank you!

 @jpbempel

